

Universal Hashing

- A malicious adversary who has learned the hash function chooses keys that all map to the same slot, giving worst-case behavior.
- Defeat the adversary using Universal Hashing
 - Use a different random hash function each time.
 - Ensure that the random hash function is independent of the keys that are actually going to be stored.
 - Ensure that the random hash function is “good” by carefully designing a class of functions to choose from.
 - Design a universal class of functions.

Universal Set of Hash Functions

- A finite collection of hash functions H that map a universe U of keys into the range $\{0, 1, \dots, m-1\}$ is “universal” if, for each pair of distinct keys, $k, l \in U$, the number of hash functions $h \in H$ for which $h(k) = h(l)$ is no more than $|H|/m$.
- The chance of a collision between two keys is the $1/m$ chance of choosing two slots randomly & independently.
- Universal hash functions give good hashing behavior.

Cost of Universal Hashing

Theorem:

Using chaining and universal hashing on key k :

- If k is not in the table T , the expected length of the list that k hashes to is $\leq \alpha$.
- If k is in the table T , the expected length of the list that k hashes to is $\leq 1 + \alpha$.

Assignment Project Exam Help

<https://powcoder.com>

Proof:

$X_{kl} = I_{\{h(k)=h(l)\}}$. $E[X_{kl}] = \Pr\{h(k)=h(l)\} \leq 1/m$.

Add WeChat powcoder

RV $Y_k =$ no. of keys other than k that hash to the same slot as k . Then,

$$Y_k = \sum_{l \in T \wedge l \neq k} X_{kl}, \text{ and } E[Y_k] = E\left[\sum_{l \in T \wedge l \neq k} X_{kl} \right] = \sum_{l \in T \wedge l \neq k} E[X_{kl}] \leq \sum_{l \in T \wedge l \neq k} \frac{1}{m}$$

If $k \notin T$, exp. length of list $= E[Y_k] \leq n/m = \alpha$.

If $k \in T$, exp. length of list $= E[Y_k] + 1 \leq (n-1)/m + 1 = 1 + \alpha - 1/m < 1 + \alpha$.

Example of Universal Hashing

When the table size m is a prime,

key x is decomposed into bytes s.t. $x = \langle x_0, \dots, x_r \rangle$,
and $a = \langle a_0, \dots, a_r \rangle$ denotes a sequence of $r+1$
elements randomly chosen from $\{0, 1, \dots, m-1\}$,

The class H defined by <https://powcoder.com>

$$H = \bigcup_a \{h_a\} \text{ with } h_a(x) = \sum_{i=0}^r a_i x_i \bmod m$$

is a universal function,

(but if some a_i is zero, h does not depend on all bytes of x and if all a_i are zero the behavior is terrible. See text for better method of universal hashing.)

Analysis on Chained-Hash-Search

- Load factor $\alpha = n/m$ = average keys per slot.
 - m – number of slots.
 - n – number of elements stored in the hash table.
- Worst-case complexity: $\Theta(n)$ + time to compute $h(k)$.
Assignment Project Exam Help
- Average depends on how h distributes keys among m slots.
- Assume
 - Simple uniform Adding WeChat powcoder
 - Any key is equally likely to hash into any of the m slots, independent of where any other key hashes to.
 - $O(1)$ time to compute $h(k)$.
- Time to search for an element with key k is $\Theta(|T[h(k)]|)$.
- Expected length of a linked list = load factor = $\alpha = n/m$.

Expected Cost of an Unsuccessful Search

Theorem 11.1:

An unsuccessful search takes expected time $\Theta(1+\alpha)$.

Proof: Assignment Project Exam Help

- Any key not already in the table is equally likely to hash to any of the m slots.
- To search unsuccessfully for any key k , need to search to the end of the list $T[h(k)]$, whose expected length is α .
- Adding the time to compute the hash function, the total time required is $\Theta(1+\alpha)$.

Expected Cost of a Successful Search

Theorem 11.2:

A successful search takes expected time $\Theta(1+\alpha)$ under simple Uniform hashing.

Proof:

- The probability that a list is searched is proportional to the number of elements it contains.
- Assume that the element being searched for is equally likely to be any of the n elements in the table.
- The number of elements examined during a successful search for an element x is 1 more than the number of elements that appear before x in x 's list.
 - These are the elements inserted *after* x was inserted.
- **Goal:**
 - Find the average, over the n elements x in the table, of how many elements were inserted into x 's list after x was inserted.

Expected Cost of a Successful Search

Theorem 11.2:

A successful search takes expected time $\Theta(1+\alpha)$ under simple Uniform hashing.

Proof (contd):

- Let x_i be the i^{th} element inserted into the table, and let $k_i = \text{key}[x_i]$.
- Define indicator random variables $X_{ij} = \{h(k_i) = h(k_j)\}$, for all i, j .
- Simple uniform hashing $\Rightarrow \Pr\{h(k_i) = h(k_j)\} = 1/m$
 $\Rightarrow E[X_{ij}] = 1/m$. (Lemma 5.1)
- Expected number of elements examined in a successful search is:

$$E\left[\frac{1}{n} \sum_{i=1}^n \left(1 + \sum_{j=i+1}^n X_{ij}\right)\right]$$

No. of elements inserted after x_i into the same slot as x_i .

Proof – Contd.

$$E\left[\frac{1}{n}\sum_{i=1}^n\left(1+\sum_{j=i+1}^n X_{ij}\right)\right]$$

$$= \frac{1}{n}\sum_{i=1}^n\left(1+\sum_{j=i+1}^n E[X_{ij}]\right) \quad (\text{linearity of expectation})$$

$$= \frac{1}{n}\sum_{i=1}^n\left(1+\sum_{j=i+1}^n \frac{1}{m}\right)$$

$$= 1 + \frac{1}{nm}\sum_{i=1}^n (n-i)$$

$$= 1 + \frac{1}{nm}\left(\sum_{i=1}^n n - \sum_{i=1}^n i\right)$$

$$= 1 + \frac{1}{nm}\left(n^2 - \frac{n(n+1)}{2}\right)$$

$$= 1 + \frac{n-1}{2m}$$

$$= 1 + \frac{\alpha}{2} - \frac{\alpha}{2n}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Expected total time for a successful search =
Time to compute hash function + Time to
search
 $= O(2 + \alpha/2 - \alpha/2n) = O(1 + \alpha)$.

Open Addressing

- An alternative to chaining for handling collisions.
- **Idea:**
 - Store all keys in the hash table itself. What can you say about α ?
 - Each slot contains either a key or NIL.
 - To *search* for key k :
 - Examine slot $h(k)$. Examining a slot is known as a **probe**.
 - If slot $h(k)$ contains key k , the search is successful. If the slot contains NIL, the search is unsuccessful.
 - There's a third possibility: **slot $h(k)$ contains a key that is not k .**
 - * Compute the index of some other slot, based on k and which probe we are on.
 - * Keep probing until we either find key k or we find a slot holding NIL.
- **Advantages:** Avoids pointers; so can use a larger table.

Probe Sequence

- Sequence of slots examined during a key search constitutes a *probe sequence*.
- Probe sequence must be a permutation of the slot numbers. **Assignment Project Exam Help**
 - We examine every slot in the table, if we have to.
 - We don't examine any slot more than once.
- The hash function is extended to:
 - $h : U \times \underbrace{\{0, 1, \dots, m-1\}}_{\text{probe number}} \rightarrow \underbrace{\{0, 1, \dots, m-1\}}_{\text{slot number}}$
- $\langle h(k,0), h(k,1), \dots, h(k,m-1) \rangle$ should be a permutation of $\langle 0, 1, \dots, m-1 \rangle$.

Ex: Linear Probing

- Example:

- $h(x) = x \bmod 13$

- $h(x,i) = (h(x) + i) \bmod 13$

- Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order

<https://powcoder.com>

Add WeChat powcoder

0	1	2	3	4	5	6	7	8	9	10	11	12



		41			18	44	59	32	22	31	73	
0	1	2	3	4	5	6	7	8	9	10	11	12

Operation Insert

- Act as though we were searching, and insert at the first NIL slot found.

- Pseudo-code for Insert:

<https://powcoder.com>

Add WeChat powcoder

```
Hash-Insert( $T, k$ )
1.  $i \leftarrow 0$ 
2. repeat  $j \leftarrow h(k, i)$ 
3.       if  $T[j] = \text{NIL}$ 
4.       then  $T[j] \leftarrow k$ 
5.       return  $j$ 
6.       else  $i \leftarrow i + 1$ 
7. until  $i = m$ 
8. error "hash table overflow"
```

Pseudo-code for Search

Hash-Search (T, k)

```
1.  $i \leftarrow 0$   
2. repeat  $j \leftarrow h(k, i)$   
3.   if  $T[j] = k$   
4.     then return  $j$   
5.      $i \leftarrow i + 1$   
6. until  $T[j] = \text{NIL}$  or  $i = m$   
7. return NIL
```

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Deletion

- Cannot just turn the slot containing the key we want to delete to contain NIL. Why?
 - We might be unable to retrieve any key k during whose insertion we had probed a slot that it occupied.
- Use a special value **DELETED** instead of NIL when marking a slot as empty during deletion.
 - **Search** should treat **DELETED** as though the slot holds a key that does not match the one being searched for.
 - **Insert** should treat **DELETED** as though the slot were empty, so that it can be reused. (So, the Hash-Insert need to be modified.)
- **Disadvantage:** Search time is no longer dependent on α .
 - Hence, chaining is more common when keys have to be deleted.

Computing Probe Sequences

- The ideal situation is *uniform hashing*:
 - Generalization of simple uniform hashing.
 - Each key is equally likely to have any of the $m!$ permutations of $\langle 0, 1, \dots, m-1 \rangle$ as its probe sequence.
- It is hard to implement true uniform hashing.
 - Approximate with techniques that at least guarantee that the probe sequence is a permutation of $\langle 0, 1, \dots, m-1 \rangle$.
- Some techniques:
 - Use *auxiliary hash functions*.
 - Linear Probing.
 - Quadratic Probing.
 - Double Hashing.
 - Can't produce all $m!$ probe sequences. (None of these can fulfill the assumption of uniform hashing.)

Linear Probing

- $h(k, i) = (h'(k) + i) \bmod m.$

key Probe number Auxiliary hash function



- The initial probe determines the entire probe sequence.
 - $T[h'(k)], T[h'(k)+1], \dots, T[m-1], T[0], T[1], \dots, T[h'(k)-1]$
 - Hence, only m distinct probe sequences are possible.
- Easy to implement, but suffers from **primary clustering**:
 - Long runs of occupied sequences build up.
 - Long runs tend to get longer, since an empty slot preceded by i full slots gets filled next with probability $(i+1)/m$.
 - Hence, average search and insertion times increase.

Ex: Linear Probing


- Example:
 - $h'(x) = x \bmod 13$
 - $h(x) = (h'(x) + i) \bmod 13$

Assignment Project Exam Help


- Insert keys 18, 41, 22, 44, 59, 32, 31, 73, in this order



Quadratic Probing

- $h(k, i) = (h'(k) + c_1 i + c_2 i^2) \bmod m$ $c_1 \neq c_2$

key Probe number Auxiliary hash function
- The initial probe position is $T[h'(k)]$, later probe positions are offset by amounts that depend on a quadratic function of the probe number i .
Assignment Project Exam Help
<https://powcoder.com>
- Must **constrain** c_1, c_2 , and m to ensure that we get a full permutation of $\langle 0, 1, \dots, m-1 \rangle$.
Add WeChat powcoder
- Can suffer from **secondary clustering**:
 - If two keys have the same initial probe position, then their probe sequences are the same. $h(k_1, 0) = h(k_2, 0)$

Double Hashing

- $h(k, i) = ((h_1(k) + i h_2(k)) \bmod m$

key Probe number Auxiliary hash functions
- Two auxiliary hash functions.
 - h_1 gives the initial probe. h_2 gives the remaining probes.
- Must have $h_2(k)$ relatively prime to m , so that the probe sequence is a full permutation of $\langle 0, 1, \dots, m-1 \rangle$.
 - Choose m to be a power of 2 and have $h_2(k)$ always return an odd number. Or,
 - Let m be prime, and have $1 < h_2(k) < m$.
- $\Theta(m^2)$ different probe sequences.
 - One for each possible combination of $h_1(k)$ and $h_2(k)$.
 - Close to the ideal uniform hashing.

0	
1	79
2	
3	
4	69
5	98
6	
7	72
8	
9	14
10	
11	50
12	

$$h(k,i) = ((h_1(k) + i h_2(k)) \bmod m)$$

$$m = 13$$

$$h_1(k) = k \bmod 13$$

$$h_2(k) = 1 + (k \bmod 11)$$

Assignment Project Exam Help

$$14 \equiv 1 \pmod{13}$$

$$14 \equiv 3 \pmod{11}$$

<https://powcoder.com>

Add WeChat powcoder

[Hashing with open addressing tool](#)

Figure 11.5 Insertion by double hashing. Here we have a hash table of size 13 with $h_1(k) = k \bmod 13$ and $h_2(k) = 1 + (k \bmod 11)$. Since $14 \equiv 1 \pmod{13}$ and $14 \equiv 3 \pmod{11}$, the key 14 is inserted into empty slot 9, after slots 1 and 5 are examined and found to be occupied.

Analysis of Open-address Hashing

- Analysis is in terms of load factor $\alpha = n/m$.
- **Assumptions:**
 - Assume that the table never completely fills, so $n < m$ and $\alpha < 1$. <https://powcoder.com>
 - Assume uniform hashing.
 - No deletion. Add WeChat powcoder
 - In a successful search, each key is equally likely to be searched for.

Expected cost of an unsuccessful

Theorem:

Given an open-address hash table with $\alpha = n/m < 1$, the expected number of probes in an unsuccessful search in an open-address hash table is at most $1/(1-\alpha)$ assuming uniform hashing.

Proof:

Assignment Project Exam Help

Every probe except the last is to an occupied slot.

Let X = # of probes in an unsuccessful search.

$X \geq i$ iff probes 1, 2, ..., $i-1$ are made to occupied slots

Let A_i = event that there is an i th probe, to an occupied slot.

$$\Pr\{X \geq i\}$$

$$= \Pr\{A_1 \cap A_2 \cap \dots \cap A_{i-1}\}.$$

$$= \Pr\{A_1\} \Pr\{A_2 | A_1\} \Pr\{A_3 | A_2 \cap A_1\} \dots \Pr\{A_{i-1} | A_1 \cap \dots \cap A_{i-2}\}$$

Proof – Contd.

$X \geq i$ iff probes 1, 2, ..., $i - 1$ are made to occupied slots

Let A_i = event that there is an i th probe, to an occupied slot.

$$\Pr\{X \geq i\}$$

$$= \Pr\{A_1 \cap A_2 \cap \dots \cap A_{i-1}\}$$

$$= \Pr\{A_1\} \Pr\{A_2 | A_1\} \Pr\{A_3 | A_2 \cap A_1\} \dots \Pr\{A_{i-1} | A_1 \cap \dots \cap A_{i-2}\}$$

<https://powcoder.com>

- $\Pr\{A_j | A_1 \cap A_2 \cap \dots \cap A_{j-1}\} = (n-j+1)/(m-j+1).$

$$\Pr\{X \geq i\} = \frac{n}{m} \cdot \frac{n-1}{m-1} \cdot \frac{n-2}{m-2} \dots \frac{n-i+2}{m-i+2}$$

$$\leq \left(\frac{n}{m}\right)^{i-1} = \alpha^{i-1}.$$

Proof – Contd.

$$\begin{aligned} E[X] &= \sum_{i=0}^{\infty} i \Pr\{X = i\} \\ &= \sum_{i=0}^{\infty} i(\Pr\{X \geq i\} - \Pr\{X \geq i+1\}) \\ &= 1 \cdot \Pr\{X \geq 1\} - 1 \cdot \Pr\{X \geq 2\} + 2 \cdot \Pr\{X \geq 2\} - 2 \cdot \Pr\{X \geq 3\} + \cdots \\ &= 1 \cdot \Pr\{X \geq 1\} + \Pr\{X \geq 2\} + \Pr\{X \geq 3\} + \cdots \quad (C.25) \\ &= \sum_{i=1}^{\infty} \Pr\{X \geq i\} \\ &\leq \sum_{i=1}^{\infty} \alpha^{i-1} = \sum_{i=0}^{\infty} \alpha^i = \frac{1}{1-\alpha} \quad (A.6) \end{aligned}$$

- If α is a constant, search takes $O(1)$ time.
- **Corollary:** Inserting an element into an open-address table takes $\leq 1/(1-\alpha)$ probes on average.

Expected cost of a successful

Theorem:

The expected number of probes in a successful search in an open-address hash table is at most $(1/\alpha) \ln (1/(1-\alpha))$.

Proof:

Assignment Project Exam Help

- A successful search for a key k follows the same probe sequence as when k was inserted.
- If k was the $(i+1)$ st key inserted, then α equaled i/m at that time.
- By the previous corollary, the expected number of probes made in a search for k is at most $1/(1-i/m) = m/(m-i)$.
- This is assuming that k is the $(i+1)$ st key. We need to average over all n keys.

Proof – Contd.

Averaging over all n keys, average # of probes is given by

$$\begin{aligned}\frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} &= \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} \\ &= \frac{1}{\alpha} (H_m - H_{m-n}) \\ &\leq \frac{1}{\alpha} \ln \frac{1}{1-\alpha}\end{aligned}$$

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Perfect Hashing

- If you know the n keys in advance, make a hash table with $O(n)$ size, and worst-case $O(1)$ lookup time!
 - Start with $O(n^2)$ size... no collisions
- <https://powcoder.com>

Thm 11.9: For a table of size $m = n^2$, if we choose h from a universal class of hash functions, we have no collisions with probability $> 1/2$.

Pf: Expected number of collisions among pairs:

$$E[X] = (n \text{ choose } 2) / n^2 < 1/2,$$

& Markov inequality says $\Pr\{X \geq t\} \leq E[X]/t$. ($t=1$)

Perfect Hashing

- If you know the n keys in advance, make a hash table with $O(n)$ size, and worst-case $O(1)$ lookup time!
- With table size n , few (collisions)²...

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Thm 11.10: For a table of size $m = n$,
if we choose h from a universal class of hash functions,

$E[\sum_j n_j^2] < 2n$, where n_j is number of keys hashing to j .

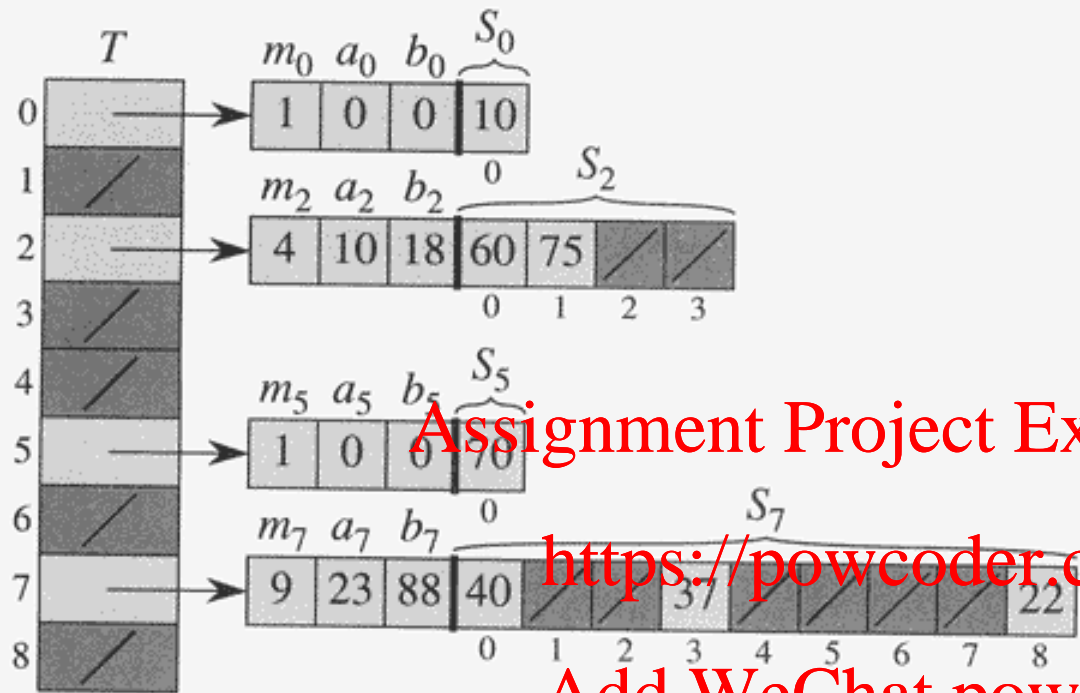
Pf: essentially the total number of collisions.

Perfect Hashing

- If you know the n keys in advance (static), make a hash table with $O(n)$ size, and worst-case $O(1)$ lookup time!
- Just use two levels of hashing:
A table of size n , then tables of size n_j^2 .



Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Figure 11.6 Using perfect hashing to store the set $K = \{10, 22, 37, 40, 60, 70, 75\}$. The outer hash function is $h(k) = ((ak + b) \bmod p) \bmod m$, where $a = 3$, $b = 42$, $p = 101$, and $m = 9$. For example, $h(75) = 2$, so key 75 hashes to slot 2 of table T . A secondary hash table S_j stores all keys hashing to slot j . The size of hash table S_j is m_j , and the associated hash function is $h_j(k) = ((a_jk + b_j) \bmod p) \bmod m_j$. Since $h_2(75) = 1$, key 75 is stored in slot 1 of secondary hash table S_2 . There are no collisions in any of the secondary hash tables, and so searching takes constant time in the worst case.

Assignment Project Exam Help
End of Chapter 11
<https://powcoder.com>

Add WeChat powcoder